

Chapter 8: Memory Management

Prof. Li-Pin Chang
CS@NYCU

Chapter 8: Memory Management


- Memory hierarchy and caching
- Address Binding
- Contiguous Memory Allocation
- Paging
- Structure of the Page Table
- Segmentation
- Example: Intel x86-32/64 & ARM32

Memory Hierarchy

Memory Hierarchy

- Main memory – the only large storage media that the CPU can access directly
 - RAM (random access), NVRAM (Non-Volatile Memory)
- Secondary storage – extension of main memory that provides large nonvolatile storage capacity
 - No random access, Magnetic disks – rigid metal or glass platters covered with magnetic recording material
 - Disk surface is logically divided into tracks, which are subdivided into sectors
 - The disk controller determines the logical interaction between the device and the computer

Memory Hierarchy (including Storage)



Storage Structure (Memory hierarchy)

- Storage systems organized in hierarchy
 - Latency/bandwidth
 - Cost/capacity
 - Volatility
- Caching – copying information into upper (faster) layer; main memory can be viewed as a last cache for secondary storage
 - Inclusive or exclusive

Caching

- Important principle, exist in many levels in a computer (hardware, operating system, applications)
 - Information in use copied from slower to faster storage temporarily
- Faster storage (cache) checked first to determine if information is there
 - If it is, information used directly from the cache (fast)
 - If not, data copied to cache and used there
- Cache design issues
 - Lookup and replacement

Performance of Various Levels of Storage


Level	1	2	3	4	5
Name	registers	cache	main memory	solid state disk	magnetic disk
Typical size	< 1 KB	< 16MB	< 64GB	< 1 TB	< 10 TB
Implementation technology	custom memory with multiple ports CMOS	on-chip or off-chip CMOS SRAM	CMOS SRAM	flash memory	magnetic disk
Access time (ns)	0.25 - 0.5	0.5 - 25	80 - 250	1ms	10ms
Bandwidth (MB/sec)	20,000 - 100,000	5,000 - 10,000	1,000 - 5,000	500	20 - 150
Managed by	compiler	hardware	operating system	operating system	operating system
Backed by	cache	main memory	disk	disk	disk or tape

CONTIGUOUS MEMORY ALLOCATION

Also known as “dynamic memory/space allocation”

Heap Management in Linux



- In Linux, a process allocates contiguous memory from its heap using `mmap(-1)` or `malloc()`
 - Heap is part of the data segment
 - If the heap is full, `malloc()` calls `brk()/sbrk()` to enlarge the data segment
- Allocated memory is returned using `free()` or `munmap()`




Memory Allocation

How to satisfy a request of size n from a list of free holes?


- **First-fit**: Allocate the *first* hole that is big enough
- **Best-fit**: Allocate the *smallest* hole that is big enough; must search entire list, unless ordered by size. Produces the smallest leftover hole.
- **Worst-fit**: Allocate the *largest* hole; must also search entire list. Produces the largest leftover hole.



Best fit



Worst fit



First fit

A20	20		108				
A15	20		15	93			
A10	20		15	10	83		
A25	20		15	10	25	58	
D20	20		15	10	25	58	
D10	20		15	10	25	58	
A8	8	12	15	10	25	58	
A30	8	12	15	10	25	30	28
D15	8	37			25	30	28
A15	8	15	22		25	30	28

First fit

D10	20	15	10		25	58		
A8	20	15	8	2	25	58		
A30	20	15	8	2	25	30	28	
D15	35		8	2	25	30	28	
A15	35		8	2	25	30	15	13

Best fit

D10	20	15	10	25	58		
A8	20	15	10	25	8	50	
A30	20	15	10	25	8	30	20
D15	45			25	8	30	20
A15	15	30		25	8	30	20

Worst fit

Buddy System

page 0 2 4 6 8 10 12 14 16

| | | | | | | | | | | | | |

Initialization

requestA (2)

requestB (1)

requestC (2)


Free A*


requestD (1)

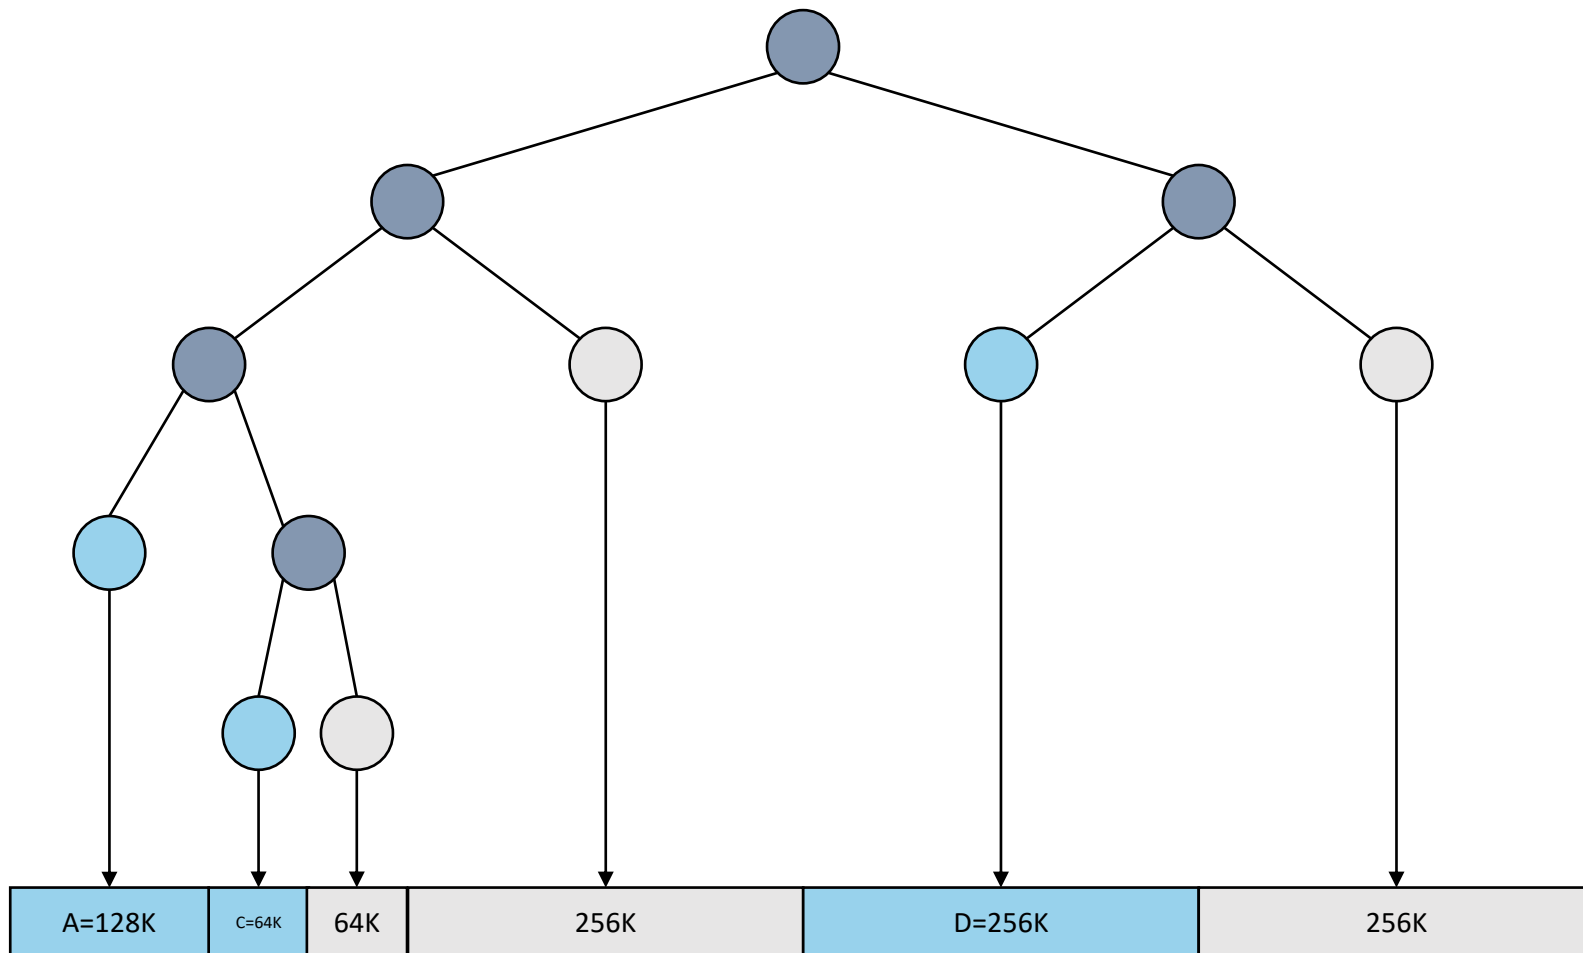
Free B

Free D

Free C







Leaf node for
allocated block



Leaf node for
unallocated block



Non-leaf node

Fragmentation

- **External Fragmentation** – fragmentation of free space; free memory sufficiently satisfies a request, but it is not contiguous
- **Internal Fragmentation** – allocated memory may be slightly larger than requested memory; this size difference is memory internal to a partition, but not being used
- Can an algorithm be subject to both?

Comparison

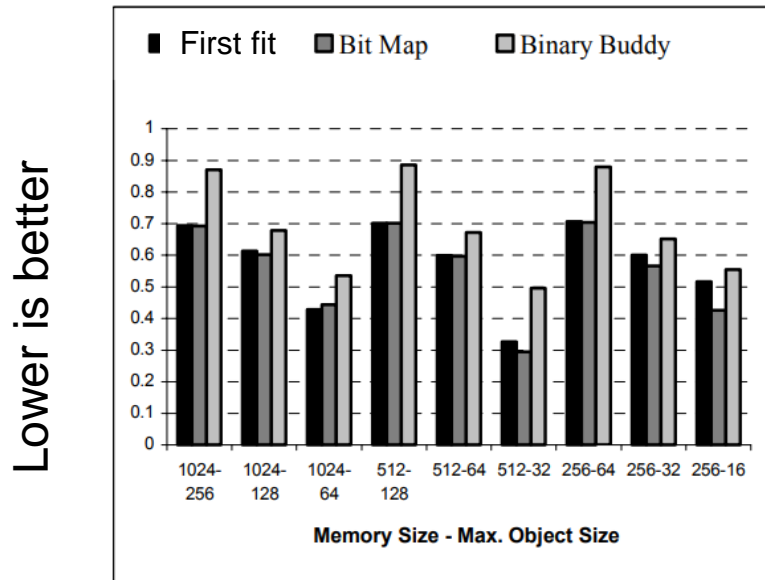


Figure 8: Total fragmentation values of techniques





Figure 9: Allocation duration of techniques (micro second)

- BF, FF, and Buddy System are practical choices
- The dynamic space allocation problem is intractable (strong NP-complete, reduced from 3-partition)

Slab Allocation

- The root cause of external fragmentation is that allocation sizes are not uniform
 - Idea: allocate objects of the same size from the same memory pool
 - Solution: slab allocation, based on the concept of size segregation
- It's part of Linux kernel memory allocation
 - Kernel objects (e.g., inode, dentry, page structure) are frequently allocated and deallocated
 - Allocation latency and memory fragmentation are crucial problems
 - Create a slab cache for each object type

Linux kernel slab cache



Benefits:

- No fragmentation
- Quick memory allocation

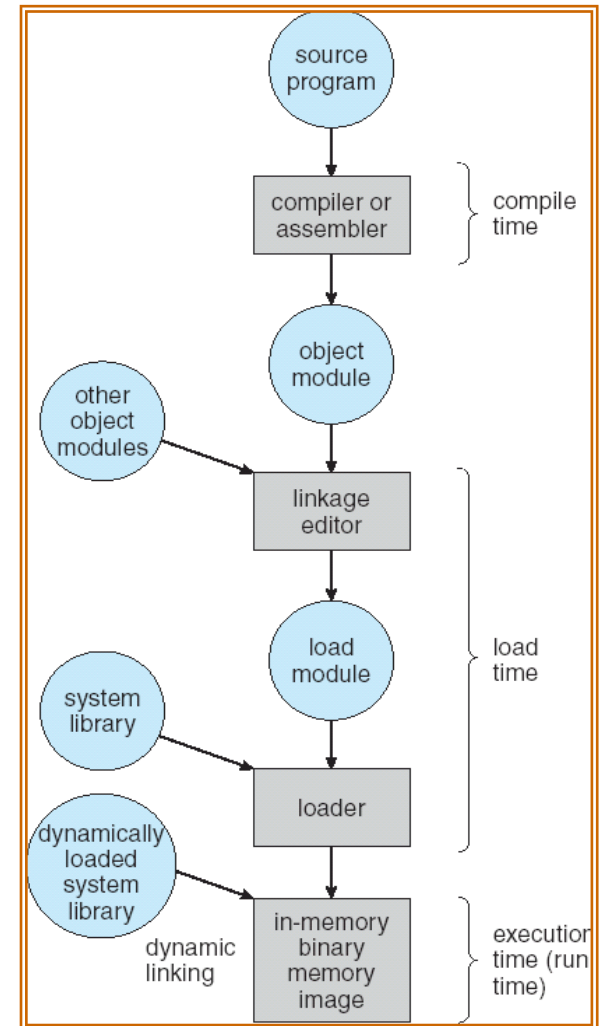
Case Study

- malloc() of glibc 2.28
 - A variation of Best Fit
 - 2^n sizes → recent chunks → small chunks → large chunks
- Linux kernel provide slab allocator and buddy system
 - Frequent allocation/deallocation of objects of the same size, use slab cache
 - Allocation/deallocation of objects of various sizes, use buddy system
- Recent enhancements to malloc()
 - Emphasis on multicore systems → thread local cache
 - [Ptmalloc](#)(GNU), [tcmalloc](#)(Google), [jemalloc](#)(Facebook)

Address Binding

Address Binding

- Assigning memory addresses to instructions and data
- Address binding can happen at different stages
 - Compile time
 - Load time
 - Execution time



Compile-Time Binding

- If memory location known a priori, absolute code can be generated; must recompile code if the starting location changes, as known as absolute addressing
- Often seen in embedded systems

Load-Time Binding

- If it is not known at compile time where the process will reside in memory, then the compiler must generate **relocatable code**. In this case, final binding is delayed until load time
- Example: PC-relative addressing (x86 assembly)

0: 66 83 f8 01

4: 74 **fa**




-6

cmp ax,0x1
je 0 <_main>

Execution-Time Binding

- Binding is delayed until run time, e.g., until a library is actually called, and the shared code can be moved to a different address during execution
- This requires hardware support, e.g., relocation registers in MMU



Logical Address vs. Physical Address

- The concept of a logical address space that is bound to a separate physical address space is central to proper memory management
 - **Logical address** – generated by the CPU; also referred to as virtual address
 - **Physical address** – address seen by the memory unit
- The user program deals with logical addresses; it never sees the real physical addresses
- Need hardware support to translate **logical addresses** into **physical addresses** during runtime
 - Paging, performed by memory management unit (MMU)


PAGING

A modern approach that separates logical (virtual) address space from physical address space



The Problem: Process Allocation



- Process are allocated to contiguous memory space
- Processes can be loaded into any contiguous and sufficiently large memory space
- As processes arrive and leave, free space will be severely fragmented



Compaction → →
To re-produce
contiguous free space



←← Hardware support for
runtime process migration:
base + limit registers

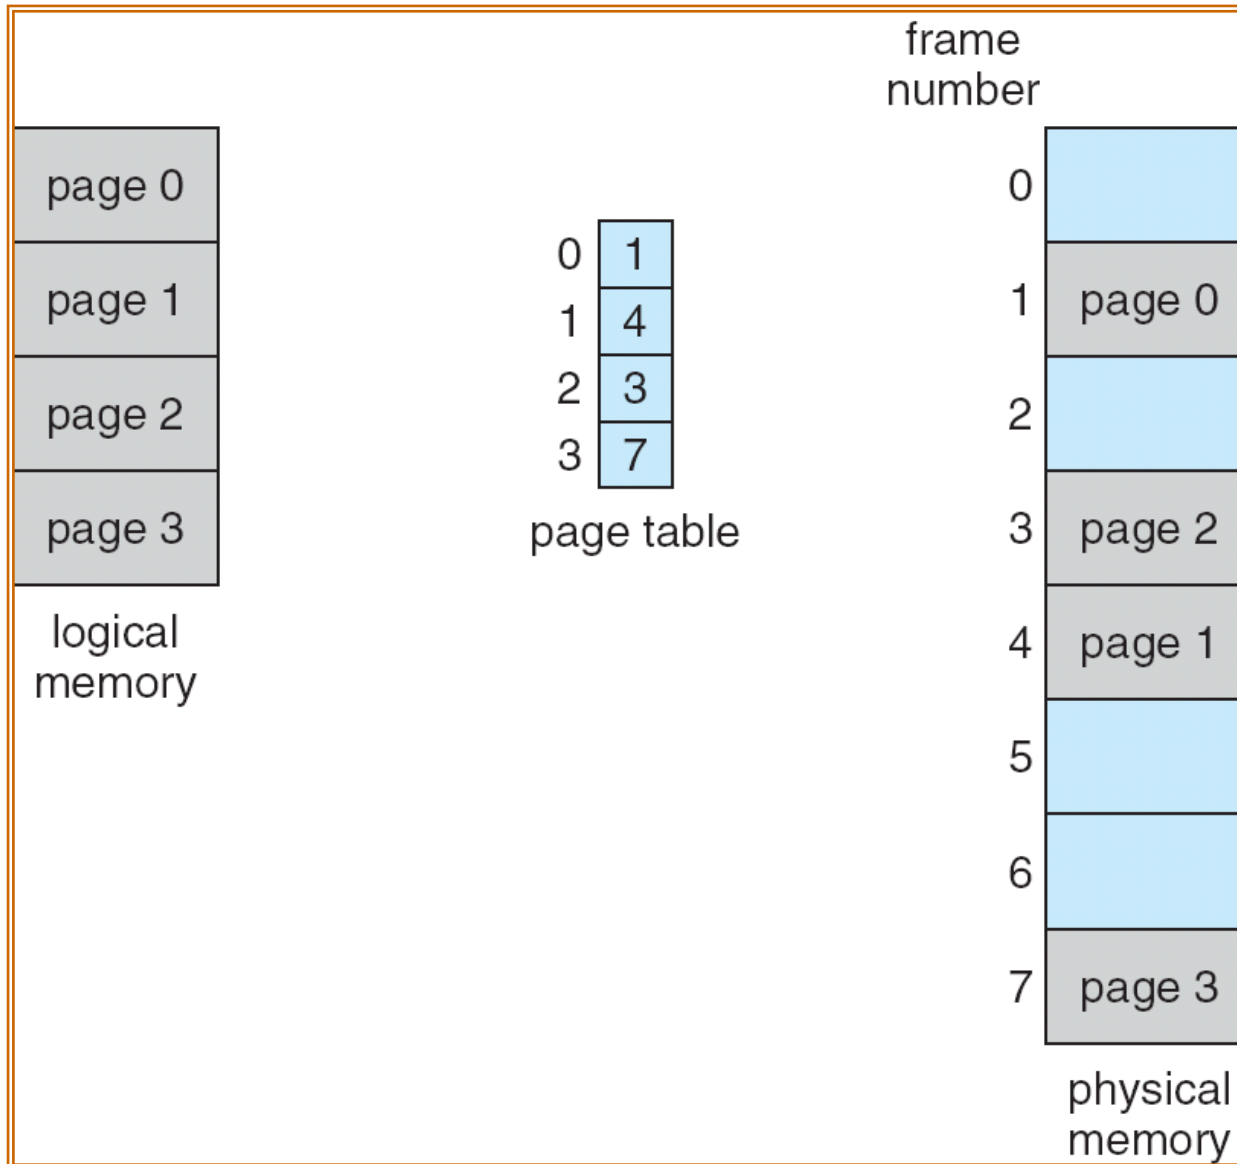
External Fragmentation

- How about compaction?
 - Migrate allocated memory chunks together to make free space contiguous
 - Will relocate programs– need execution time binding
 - Occasionally slows down the system: too costly!
- A better approach: What if physically fragmented free spaces can logically be treated contiguous?


Paging

- Divide **physical** memory into fixed-sized blocks called **frames** (size is a power of 2, typically 4KB)
- Divide **logical** memory into blocks of same size called **pages** (page size = frame size)
- Keep track of all free frames
- To run a program of size n pages, need to find n free frames and load program
- Set up a page table to translate logical to physical addresses


Paging Model of logical and physical memory



Paging Hardware



Paging Example




Q: Where is the page table? In the CPU? In memory?

Q: how many bits are required to represent


1. A logical address
2. The page number
3. A physical address
4. The frame number
5. The displacement
6. A page-table entry?

Memory-Management Unit (MMU)

- A hardware component in the CPU that translates logical addresses into physical addresses
- Enabling paging and/or segmentation




[Project Denver 64-bit CPU core](#)



[Intel i7-5960X](#)

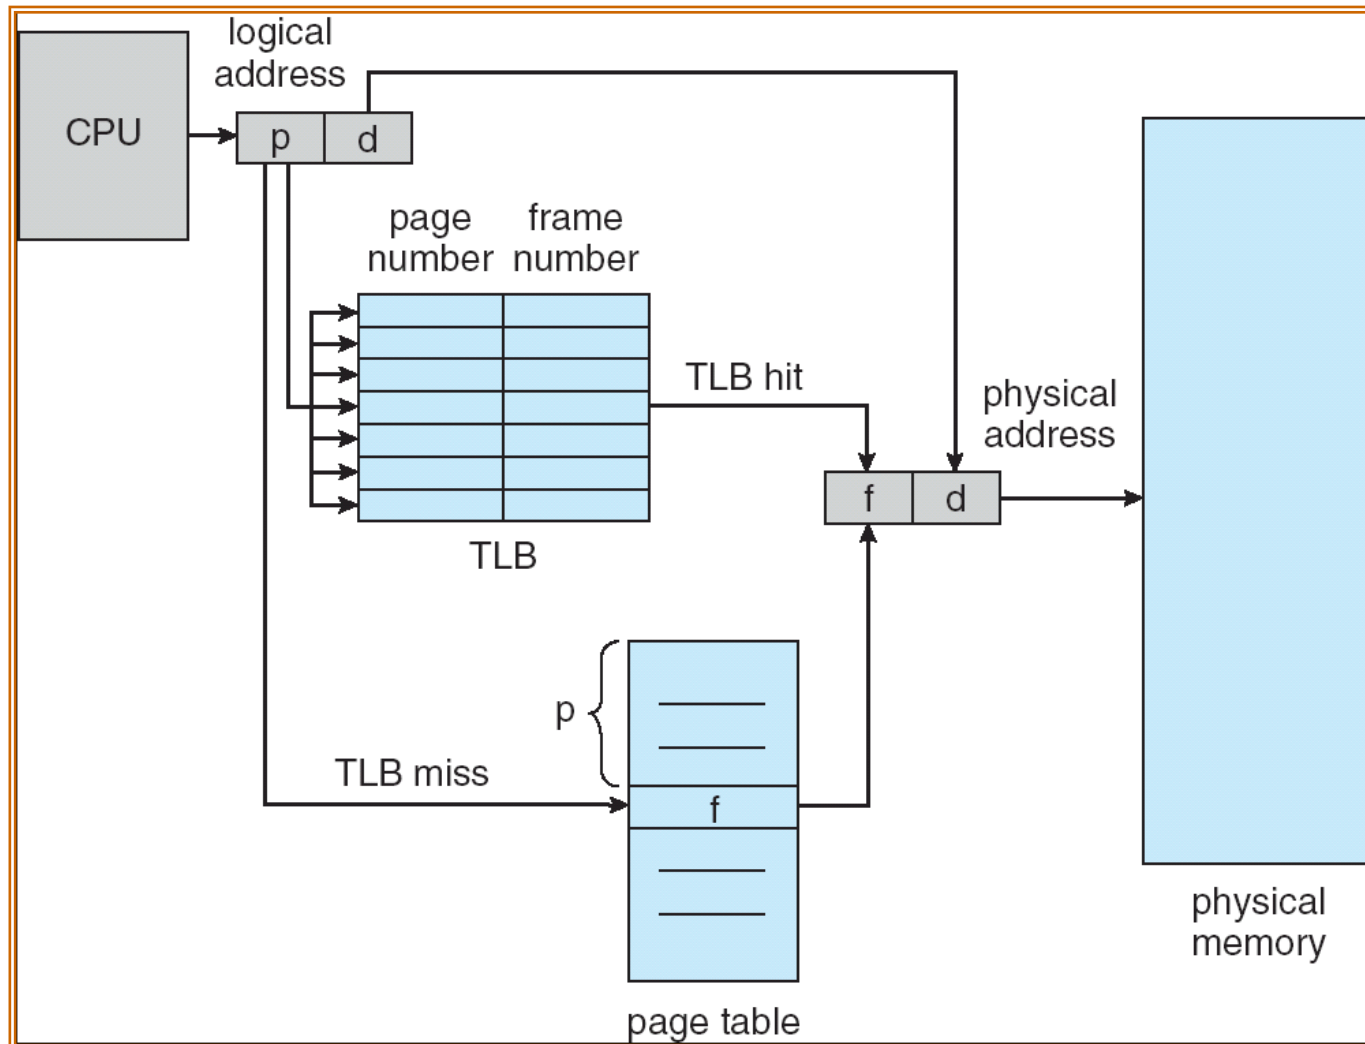
Free Frames



Implementation of Page Table

- Page table is kept in **main memory (!)**
- **A table per process**
 - Page-table base register (PTBR) points to the page table
 - Page-table length register (PRLR) indicates size of the page table
- In this scheme every data/instruction access requires **two** memory accesses. One for the page table and one for the data/instruction.
- The two memory access problem can be solved by the use of a special fast-lookup hardware cache called **associative memory** or **translation look-aside buffers (TLBs)**


Paging Hardware With TLB



All the steps are handled by HW. Data/inst → cache; page table → TLB

TLB hit ratio vs. Locality of Reference

- TLB is small, usually holds 64~1024 entries
 - A replacement policy should be used.
 - E.g., random policy or LRU
 - Page references have **temporal** localities and **spatial** localities
- Important entries can be wired down (nailed)
 - E.g., kernel code



Effective Access Time

- Associative Lookup = ε time unit
- Assume memory cycle time is 1 microsecond
- Hit ratio α – percentage of times that a page number is found in the associative registers; ratio related to number of associative registers
- **Effective Access Time (EAT)**

$$\begin{aligned} \text{EAT} &= (1 + \varepsilon) \alpha + (2 + \varepsilon)(1 - \alpha) \\ &= 2 + \varepsilon - \alpha \end{aligned}$$

Let the TLB hit ratio be 98%

20ns to lookup the TLB

100 ns to access memory

TLB hit: 20+100

TLB miss: 20 + 100 (page table) + 100 (target address)

$EAT = 0.98 * 120 + 0.02 * 220 = 122 \text{ ns}$

STRUCTURE OF THE PAGE TABLE

Structure of Page Table


- Hierarchical Paging
- Hashed Page Tables
- Inverted Page Tables

Hierarchical Page Tables

- A per-process page table could be very large and sparse
- Allocating **Large** and **contiguous** page tables for processes is inconvenient and may under-utilize memory space (of page tables)
- Breaking up the logical address space into multiple page tables
 - Page table **can be divided into pieces**
 - Page table pieces are **allocated on demand**
- A simple technique is a two-level page table
 - Example: Intel 80386

Two-Level Paging Example

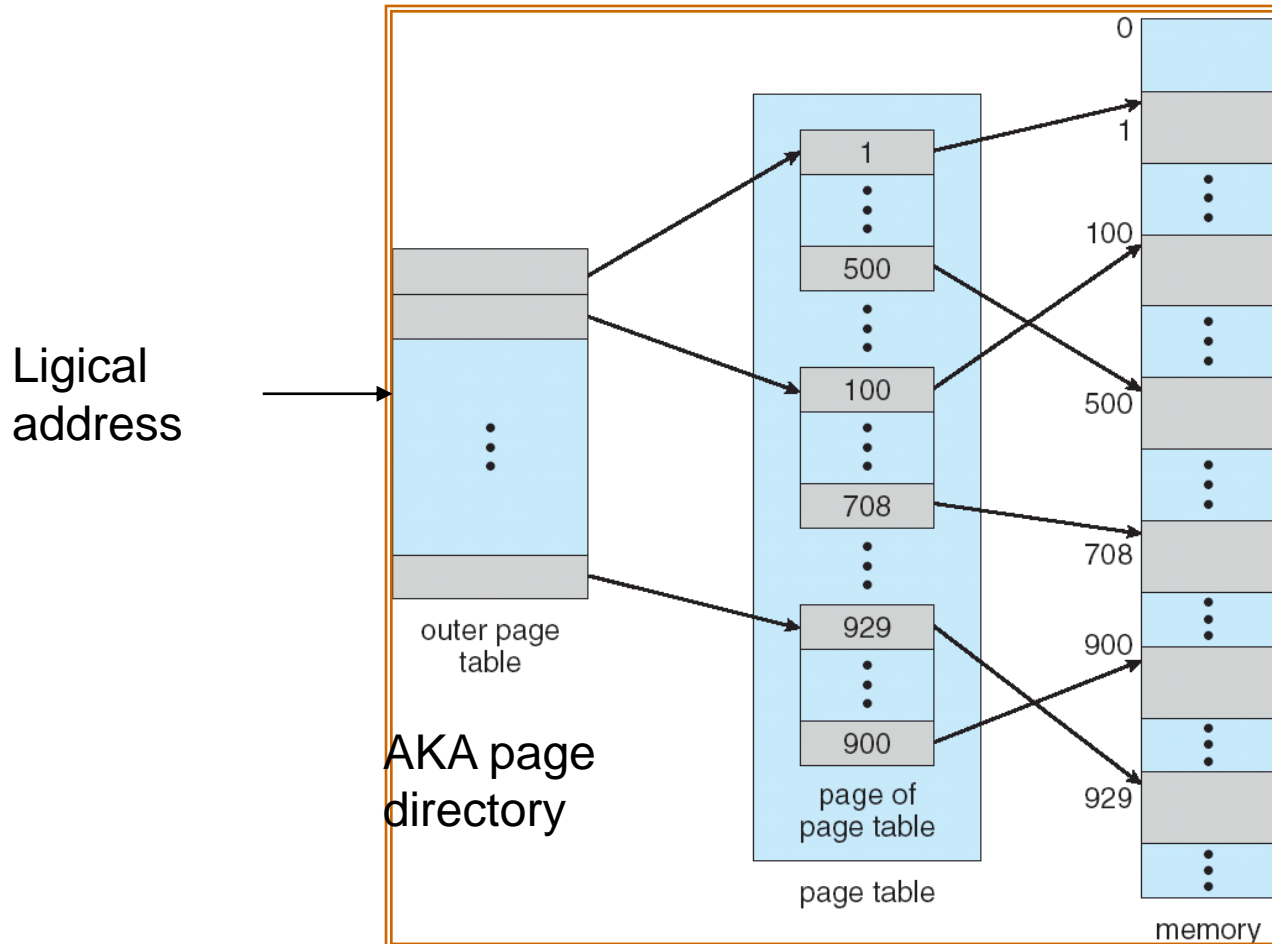
- A logical address (on 32-bit machine with 4K page size) is divided into:
 - a page number consisting of 20 bits
 - a page offset consisting of 12 bits
- Since the page table is paged, the page number is further divided into:
 - a 10-bit outer-page number
 - a 10-bit displacement of the outer page
- Thus, a logical address is as follows:



where p_i is an index into the outer page table, and p_2 is the displacement within the page of the outer page table

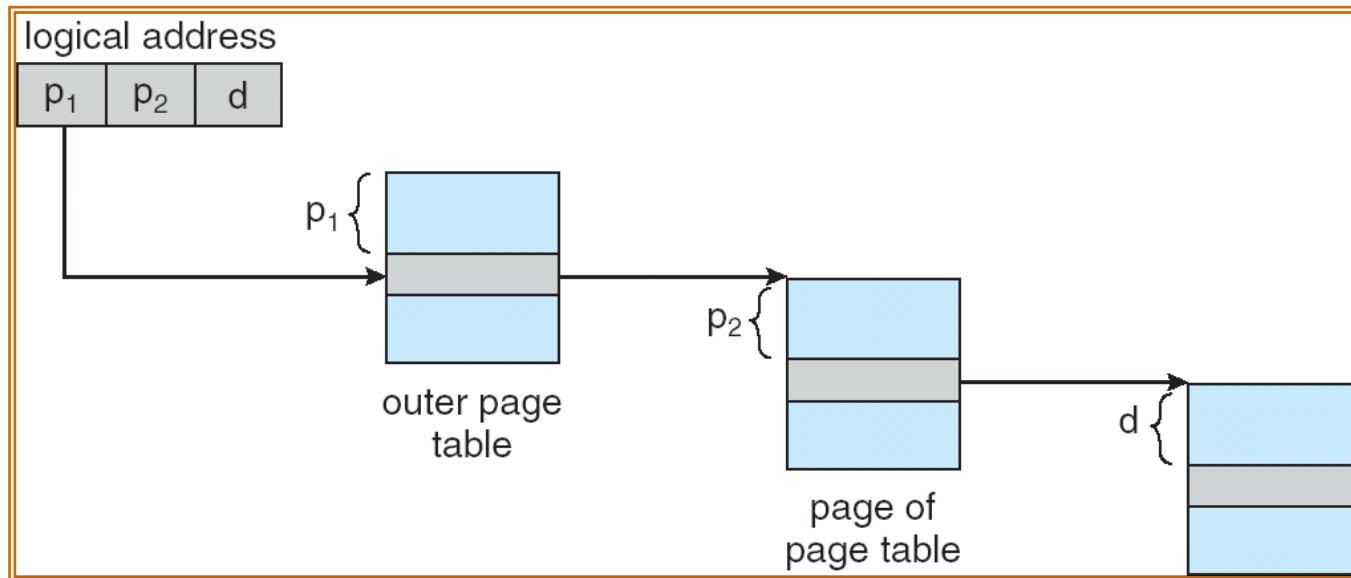
Two-Level Page-Table Scheme

[p1][p2][d]



Address-Translation Scheme


- Address-translation scheme for a two-level 32-bit paging architecture



Pros: page tables need not to be contiguous, and need not all present in memory
Cons: multiple memory accesses on TLB miss. 7-level paging in UltraSparc 64

Think about it...


- Logical address or physical address?



Hashed Page Tables

- Common in address spaces > 32 bits
- Replace the radix-based multi-level table with a hash table
 - With a huge and sparse virtual address space, a TLB miss always cause many references to the multilevel page table
 - For large, sparse virtual address spaces, an adequately-sized hash table may require one or two accesses only
- The page number is hashed into a page table. This page table contains a chain of elements hashing to the same location
- Page numbers are compared in this chain searching for a match. If a match is found, the corresponding physical frame is extracted

Hashed Page Table



How many extra memory references are needed on TLB miss is related to the quality of the hash function and the hash table size

Hashed Page Table

- (Forward) page tables ($p \rightarrow f$) are commonly implemented using multi-level tables, not hash tables
- Inverted page tables (to be shown) are, however, commonly implemented using hash tables


Inverted Page Table

- Each process has its own (forward) page table, and assume that a page table is of 4MB and there are 2000 processes running in the system==> total memory requirement of page tables is ~8GB
- The size of an inverted page table is bounded by the size of physical memory and is irrelevant to # of processes

Inverted Page Table

- One global page table shared by all processes
- One entry for each real page (frame) of memory
 - Entry index number is the frame number
 - Each table entry stores a page number plus the process ID that owns the page
- On memory reference, find the table entry that stores the current process ID and the referenced page number
 - Use hash table to limit the search to one — or at most a few — page-table entries
- Example: PowerPC 603


Inverted Page Table Architecture



“search” part is usually implemented using a hash function; pages sharing the same hash value can be chained on a link list

Think about it...

- Why **pid** is necessary in inverted page tables?



Design Considerations of Paging

- Increased memory references
 - Solution: Use TLB to eliminate PT references

Forward page table designs

- Linear table
- Multilevel page tables
 - Smaller tables, allocate on demand
- Hash page tables
 - Small hash table with overflow handling (linked lists)

Inverted page tables


- Implemented using hash table
- Table size is bounded by the physical memory size and is independent of the total number of active processes

SEGMENTATION


Segmentation

- Memory-management scheme that supports user view of memory
- A program is a collection of segments. A segment is a logical unit such as:


main program,
procedure,
function,
method,
object,
local variables, global variables,
common block,
stack,
symbol table, arrays



Logical View of Segmentation



user space



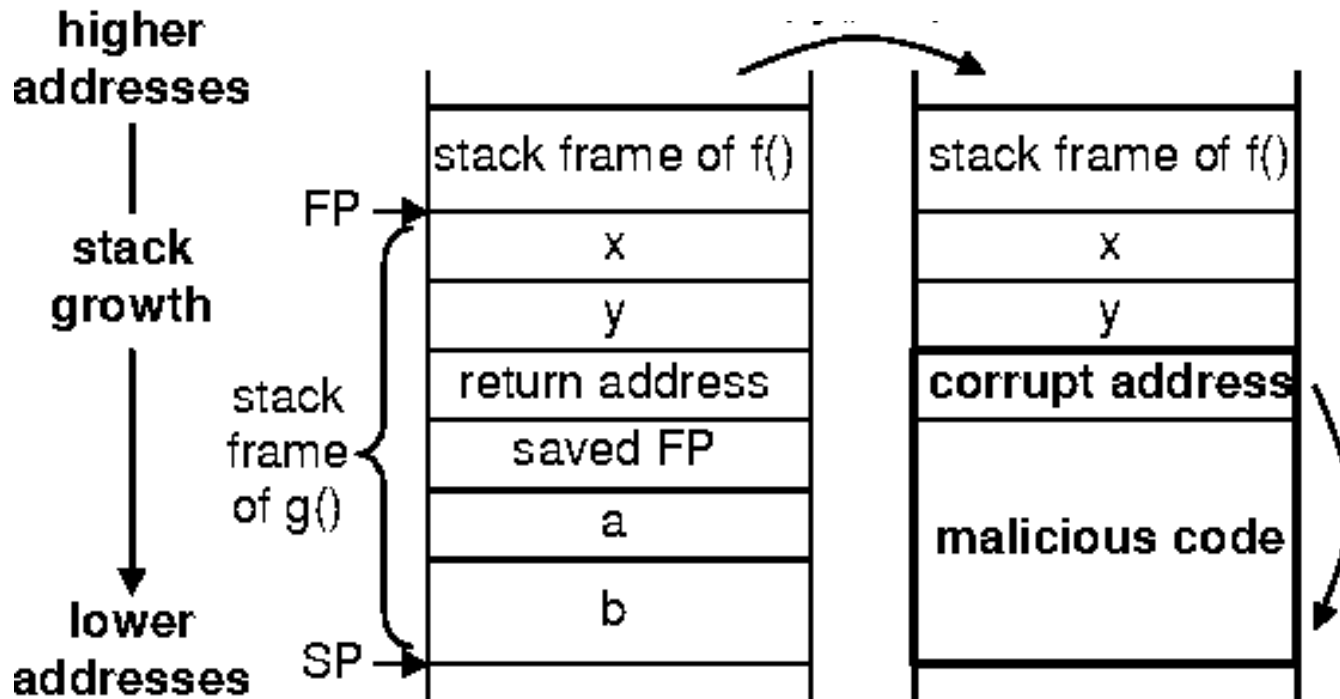
physical memory space

User programs do not know where in physical memory a segment is placed.

Common Segments and Their Permissions

- Stack segment: rw-
 - Stack frame for function calls, including local variables, function arguments, return address, CPU flags
- Data segment:
 - rw-: bss (block started by symbol, uninitialized global variables)
 - rw-: Data (initialized global variables)
 - r--: Data RO (constant literals, strings)
 - Rw-: Heap (for dynamic memory allocation)
- Code segment
 - r-x: Executable binaries (r is for copy-on-write)

Think about it: Stack Segment



Stack Overrun Attack

Think about it: Data Segment

- What is the result of executing the program below?

```
#include<stdio.h>
#include<stdlib.h>

int main()
{
    char *p = "hello world";
    p[0] = 'H';
}
```


Segmentation Hardware

- Logical address consists of a two tuple:
 <segment-number, offset>,
- **Segment table** – maps two-dimensional physical addresses; each table entry has:
 - **base** – contains the starting physical address where the segments reside in memory
 - **limit** – specifies the length of the segment
- Segment-table base register (STBR) points to the segment table's location in memory
- Segment-table length register (STLR) indicates number of segments used by a program;
 segment number s is legal if $s < \text{STLR}$

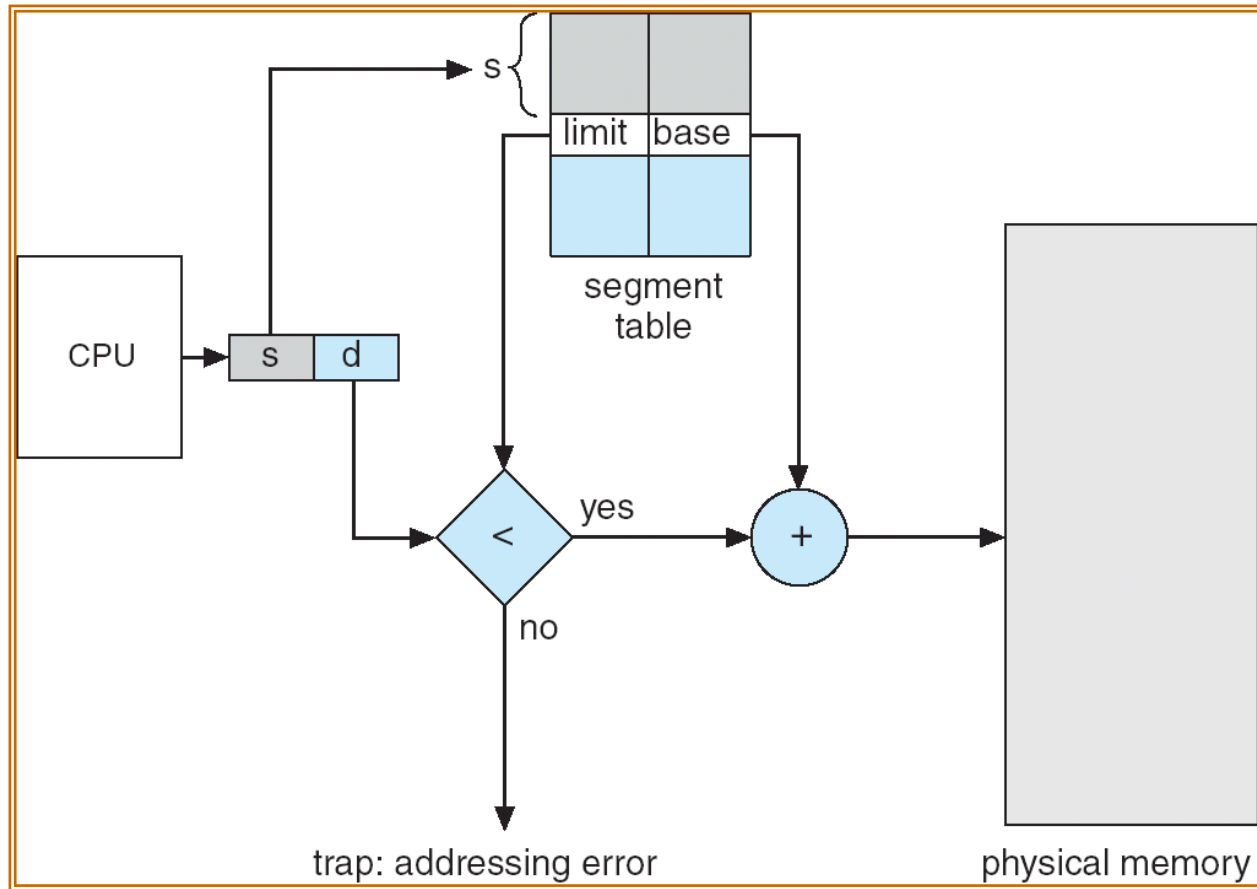
pmap example

- RSS = resident set size, how much physical memory in use (KB)
- Dirty = modified memory size (KB)
- rwx = permissions

```
$ pmap -x 78123
export "PS1=$ "
pmap -x 78123


78123:  bash
Address      Kbytes      RSS      Dirty Mode  Mapping
0000558e46119000    188      188        0 r---- bash
0000558e46148000    892      892        0 r-x-- bash
0000558e46227000    232      164        0 r---- bash
0000558e46262000     16       16       16 r---- bash
0000558e46266000     36       36       36 rw--- bash
0000558e4626f000     44       28       28 rw--- [ anon ]
0000558e46781000    264      264      264 rw--- [ anon ]
00007f011c488000     12        8        8 rw--- [ anon ]
00007f011c48b000    160      160        0 r---- libc.so.6
00007f011c4b3000   1620     1268        0 r-x-- libc.so.6
00007f011c648000    352      160        0 r---- libc.so.6
00007f011c6a0000     16       16       16 r---- libc.so.6
00007f011c6a4000      8        8        8 rw--- libc.so.6
00007f011c6a6000     52       24       24 rw--- [ anon ]
00007f011c6b3000     56       56        0 r---- libtinfo.so.6.3
00007f011c6c1000     68       64        0 r-x-- libtinfo.so.6.3
00007f011c6d2000     56       56        0 r---- libtinfo.so.6.3
00007f011c6e0000     16       16       16 r---- libtinfo.so.6.3
00007f011c6e4000      4        4        4 rw--- libtinfo.so.6.3
00007f011c6fb000      8        8        8 rw--- [ anon ]
00007f011c6fd000      8        8        0 r---- ld-linux-x86-64.so.2
00007f011c6ff000    168      168        0 r-x-- ld-linux-x86-64.so.2
00007f011c729000     44       44        0 r---- ld-linux-x86-64.so.2
00007f011c735000      8        8        8 r---- ld-linux-x86-64.so.2
00007f011c737000      8        8        8 rw--- ld-linux-x86-64.so.2
00007ffef3b73000    132     104      104 rw--- [ stack ]
00007ffef3bf8000     16        0        0 r---- [ anon ]
00007ffef3bfc000      8        4        0 r-x-- [ anon ]
fffffffffff60000      4        0        0 --x-- [ anon ]
-----
total kB          4496     3780     548
export "PS1=$ "
```

Segmentation Hardware



To this point, similar to relocation registers

Example of Segmentation



Segment Protection

- Segments storing execution code normally do not allow modification
 - Allow read, execute but not write
 - Avoid self-modifying malicious code, such as polymorphic viruses
- Stack segments normally do not allow execution
 - Allow read, write but not execute
 - Avoid malicious code injection

Hardware Support for Segmentation

- Provided by MMU
- Segment Limit
 - Segment base register
 - Segment limit register
- Segment Protection
 - With each entry in segment table associate:
 - **Permissions** of read, write, and/or execute
 - **Access privileges** of user mode or kernel mode

However.....

Paging vs. Segmentation, and the Reality...

- The functionality of paging and segmentation are, however, highly overlapped, e.g., paging hardware also supports read-write protection
- Segmentation implementation varies a lot between different CPUs
- Due to the two reasons, for better portability, **modern operating systems tend to use minimal segmentation**
- In fact, in Linux, the concept “segmentation” (aka sectioning or mapping) is implemented with paging, ***not*** segmentation!

Linux Segmentation on Intel 80386

- Segmentation and Paging are somewhat redundant
 - RW protection → supported by both
 - Executable → supported by both*
 - User/kernel Privilege → supported by both
- RISC architectures often have limited support for segmentation
- Therefore, Linux use segmentation only when required by x86 architecture, mainly on purposes for context-switching

**page executable bit is supported in x86-64 only, not x86-32. For x86-32, page execution permission is implemented through the split (I/D) TLB mechanism, see <https://pax.grsecurity.net/docs/pageexec.txt>

Linux Segmentation on Intel 80386


- Uses minimal segmentation to keep memory management implementation more portable
- Uses 6 global segments: (80386 has many)
 - Shared by all processes:
 - Kernel code, Kernel data
 - User code, User data
 - The default LDT (usually not used)
 - Per-core
 - Task-state (TSS, used to switch from user mode to kernel mode)
- Uses 2 protection levels: (80386 has 4)
 - Kernel mode / User mode

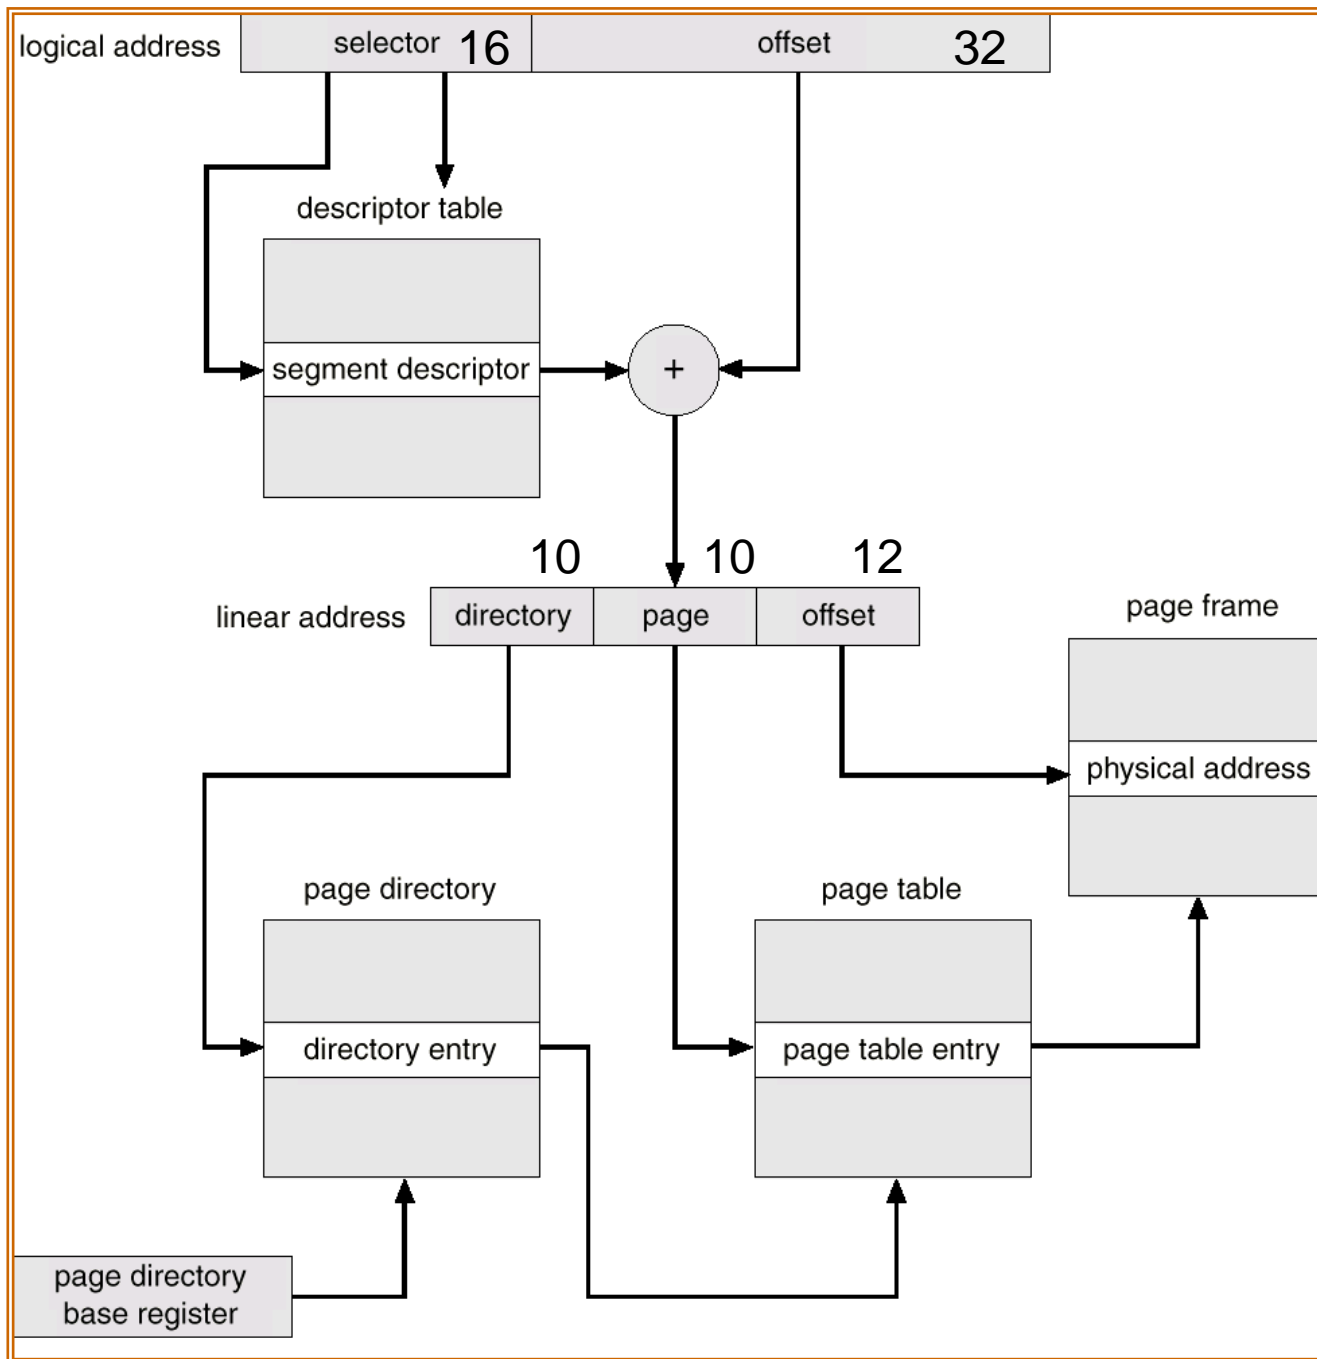
EXAMPLE: THE INTEL 80386+

Although segmentation is minimally used, let's see how segmentation and paging are put together in real hardware...

Segmented Paging


- Intel 80386
 - Segmentation
 - 2-level paging






Intel x86-64

- Current generation Intel x86 architecture
- 64 bits is ginormous (> 16 exabytes)
- In practice only implement 48 bit addressing
 - Page sizes of 4 KB, 2 MB, 1 GB
 - Four levels of paging hierarchy
- Can also use PAE so virtual addresses are 48 bits and physical addresses are 52 bits



Example: ARM Architecture

- Dominant mobile platform chip (Apple iOS and Google Android devices for example)
- Modern, energy efficient, 32-bit CPU
- 4 KB and 16 KB pages
- 1 MB and 16 MB pages (termed sections)
- One-level paging for sections, two-level for smaller pages
- Two levels of TLBs
 - Outer level has two micro TLBs (one data, one instruction)
 - Inner is single main TLB
 - First inner is checked, on miss outer are checked, and on miss page table walk performed by CPU



End of Chapter 8

Review Questions

1. It is a common issue that in tiny embedded devices, long-term dynamic memory allocation will fail even if the total available memory is sufficient. Discuss the reason why, and propose a workaround
2. In early 32-bit processors, the TLB has 128 entries typically. Discuss why such a small TLB achieves a reasonably high TLB hit rate
3. In multilevel paging, do the intermediate page tables use physical pointers or virtual pointers to the next-level page tables?
4. Why inverted page tables are often implemented as a hash table?
5. How do you implement shared memory using inverted page tables?
6. Discuss why do modern operating systems use minimal segmentation but rely on paging for memory protection
7. Based on the x86-32 architecture, calculate how many bits (in minimal) are used by 1) a page directory entry, 2) a page table entry, 3) a physical address, and 4) a virtual address
8. X86-32 paging does not support executable bit for pages. Survey how does it implement the executable permission through TLB